



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbit.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation



MINI PROJECT REPORT

on

"RN-DPE-RDHEI: Reversible Data Hiding in Encrypted Images Using Residual CNN Predictor, BRBCC Compression, and Chaos-Based Difference-Preserving Encryption"

Submitted in partial fulfilment for the completion of

BE – VI Semester

in

INFORMATION TECHNOLOGY

By

Bassa Srilakshmi (160123737005)

Under the guidance of

Mr. N. Shiva Kumar

Assistant Professor,

Dept. of IT, CBIT

DEPARTMENT OF INFORMATION TECHNOLOGY

CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY

(Autonomous)

(Affiliated to Osmania University; Accredited by NBA(AICTE) and NAAC(UGC), ISO Certified
9001:2015)

GANDIPET, HYDERABAD – 500 075

Website: www.cbit.ac.in



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**
An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in

Approved by Affiliated to UGC Autonomous
12 Programs Accredited by Grade A++ in All India Ranking 151-200 Band Certified by ISO 21001 : 2018

COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

2025–2026



CERTIFICATE

This is to certify that the project report entitled "RN-DPE-RDHEI: Reversible Data Hiding in Encrypted Images Using Residual CNN Predictor, BRBCC Compression, and Chaos-Based Difference-Preserving Encryption" is a bonafide work carried out by the following student of B.E. (Information Technology) VI Semester, Chaitanya Bharathi Institute of Technology, Hyderabad, in partial fulfillment of the requirements for the award of the degree of Bachelor of Engineering during the academic year 2025–2026. The work has been carried out under our guidance and supervision.

S.No.	Student Name	Roll Number	Signature
1.	Bassa Srilakshmi	160123737005	

Project Guide	Head of the Department
Mr. N. Shiva Kumar	Dr. M. Venu Gopala Chari
Assistant Professor	Professor & Head
Dept. of IT, CBIT	Dept. of IT, CBIT

External Examiner Signature: _____



DECLARATION BY STUDENT

I, the undersigned student of B.E. (Information Technology) VI Semester, Chaitanya Bharathi Institute of Technology (Autonomous), Hyderabad, hereby declare that the project report entitled "RN-DPE-RDHEI: Reversible Data Hiding in Encrypted Images Using Residual CNN Predictor, BRBCC Compression, and Chaos-Based Difference-Preserving Encryption" submitted in partial fulfillment of the requirements for the award of the degree of Bachelor of Engineering in Information Technology is a record of original work done by me during the academic year 2025–2026 under the supervision and guidance of Mr. N. Shiva Kumar, Assistant Professor, Department of Information Technology, CBIT, Hyderabad.

I further declare that:

- This project has NOT been submitted previously, in part or in full, to this University or any other University or Institution for the award of any degree or diploma.
- The work presented in this report is entirely my own and has been carried out with honesty and integrity.
- All sources of information, data, and references used in this report have been duly cited and acknowledged in compliance with academic norms.
- The similarity index of this report as verified by Turnitin is within the permissible limit of 20% for the project report and below 10% for the attached research paper.

Bassa Srilakshmi Roll No.: 160123737005 Signature: _____

Date: _____

Place: Hyderabad



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

ACKNOWLEDGEMENTS

I consider it a privilege to express my heartfelt gratitude to all those who extended their support and guidance in the successful completion of this project.

I extend my sincere and deepest gratitude to my project guide, Mr. N. Shiva Kumar, Assistant Professor, Department of Information Technology, CBIT, Hyderabad, for his invaluable guidance, constant encouragement, and constructive suggestions throughout the course of this project. His deep expertise in image security and data hiding research provided the technical direction that shaped every stage of this work. His insightful feedback and patient mentoring have been the foundation of this report.

I express my profound thanks to Dr. M. Venu Gopala Chari, Professor and Head, Department of Information Technology, CBIT, Hyderabad, for his encouragement, administrative support, and for providing all necessary resources and facilities that contributed to the successful execution of this project. His vision for research excellence in the department has been a constant source of motivation throughout my academic journey.

I am deeply indebted to Dr. C. V. Narasimhulu, Principal, Chaitanya Bharathi Institute of Technology, Hyderabad, for his visionary leadership, unwavering support, and for fostering a research-oriented environment that enabled me to undertake and complete this work.

I also thank all the faculty members and staff of the Department of Information Technology for their kind cooperation, timely advice, and moral support extended to me during the project period. The computational resources provided through the Kaggle platform with NVIDIA T4 GPU were essential for training the deep learning components of this system.



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**
An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in

Approved by
Affiliated to
UGC Autonomous
12 Programs
Accredited by
Grade A++ in
All India Ranking
151-200 Band
Certified by
ISO
21001 : 2018

COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

Finally, I express my deep sense of gratitude to my family and friends for their continuous motivation, patience, and unconditional support throughout my academic journey. Their encouragement during long hours of experimentation and writing was truly invaluable.

Bassa Srilakshmi



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

ABSTRACT

This report presents RN-DPE-RDHEI (Residual Network-based Difference-Preserving Encryption for Reversible Data Hiding in Encrypted Images), a unified framework that simultaneously achieves state-of-the-art embedding capacity, near-zero prediction accuracy, and strong cryptographic security. All existing Reserve Room Before Encryption (RRBE) approaches depend on hand-crafted algebraic predictors such as MED, GAP, and ISGAP, which impose a hard performance ceiling on embedding capacity. The proposed system introduces the first learned residual Convolutional Neural Network predictor in the RDHEI literature — a compact four-layer model trained to correct Median Edge Detector residuals — reducing prediction Mean Absolute Error from 1.2851 to 0.0516 pixels, a 96% reduction. Block-level Reservation via Bit-plane Compression Coding (BRBCC) is applied to the CNN-predicted error image, freeing six bit-planes per usable 4×4 block and raising embedding capacity from 3.9381 BPP (BRBCC+MED, prior state-of-the-art) to 5.9578 BPP — a 51% improvement representing the highest embedding capacity reported on BOSSBase 1.01 in the published literature. A Difference-Preserving Cipher seeded by ECDH X25519 with HKDF-SHA256 key derivation encrypts the reserved image while preserving BRBCC block validity in the ciphertext domain. Dual HMAC-SHA256 and ECDSA Ed25519 authentication — the first such combination in any RRBE scheme — provide data integrity and sender non-repudiation. Experiments on 1,000 BOSSBase 1.01 test images confirm $\text{PSNR} = \infty$, $\text{SSIM} = 1.0$, $\text{BPP} = 5.9578$, $\text{Entropy} = 7.9959$, $\text{NPCR} = 99.61\%$, $\text{UACI} = 34.29\%$, $\text{MAE} = 0.0516$ pixels, with zero failures and perfect secret recovery across all test images.

Keywords: Reversible Data Hiding, Encrypted Images, CNN Residual Predictor, BRBCC Compression, ECDH X25519, Chaos Encryption, Difference-Preserving Cipher, HMAC-SHA256, ECDSA Ed25519.



TABLE OF CONTENTS

CHAPTER / SECTION TITLE	PAGE
Certificate	ii
Declaration	iii
Acknowledgements	iv
Abstract	v
Table of Contents	vi
List of Figures	vii
List of Tables	viii
List of Abbreviations	ix
CHAPTER 1. INTRODUCTION	1
1.1 Background and Motivation	1
1.2 Problem Statement	3
1.3 Objectives	4
1.4 Scope of Project	5
1.5 Organisation	6
CHAPTER 2. LITERATURE SURVEY	7
2.1 Foundational RDHEI Frameworks	7
2.2 Predictor-Based Methods	9
2.3 Compressor-Based Methods	11
2.4 Security Extensions	12
2.5 Gap Analysis	13
CHAPTER 3. PROPOSED METHODOLOGY	15
3.1 System Overview	15
3.2 Key Exchange (ECDH X25519)	16
3.3 MED + CNN Prediction (ACNNP)	18
3.4 BRBCC Space Reservation	20
3.5 DPC Cipher (Perm + ZeroFree + XOR)	22
3.6 Dual Authentication	24
3.7 Three-Party Protocol	25



CHAPTER 4. IMPLEMENTATION	27
4.1 Technology Stack	27
4.2 CNN Architecture & Training	28
4.3 Key Management Code	30
4.4 BRBCC & Embedding Code	32
4.5 Cipher & Authentication Code	34
CHAPTER 5. TESTING AND RESULTS	36
5.1 Experimental Setup	36
5.2 Prediction Accuracy	37
5.3 Embedding Capacity	38
5.4 Full Pipeline Evaluation	39
5.5 Security Analysis	41
5.6 State-of-the-Art Comparison	43
5.7 Functional Test Cases	44
CHAPTER 6. CONCLUSION AND FUTURE SCOPE	46
6.1 Conclusion	46
6.2 Limitations	47
6.3 Future Scope	48
CHAPTER 7. VISIBLE OUTPUT (RESEARCH PAPER)	50
BIBLIOGRAPHY	52
APPENDIX	56



LIST OF FIGURES

Figure No.	Figure Title	Page
Fig. 1.1	RDHEI Three-Party System Overview	2
Fig. 1.2	RRBE vs VRAE Architecture Comparison	4
Fig. 3.1	Complete RN-DPE-RDHEI 11-Stage Pipeline	15
Fig. 3.2	ECDH X25519 Key Exchange and HKDF Derivation	17
Fig. 3.3	ResidualACNNP CNN Architecture (4 Layers, 148,353 params)	19
Fig. 3.4	BRBCC Block Classification: Class 0, 1, and 2	21
Fig. 3.5	Chaos Sequence Generation: LSS and Lorenz ODE	23
Fig. 3.6	DPC Cipher: Permutation + ZeroFree + XOR Steps	23
Fig. 3.7	Dual Authentication: HMAC-SHA256 + ECDSA Ed25519	24
Fig. 4.1	CNN Training and Validation Loss Curves (50 Epochs)	29
Fig. 4.2	Free-Plane Scan Pattern for Embedding and Extraction	33
Fig. 5.1	MAE Comparison: All RDHEI Predictors on BOSSBase 1.01	37
Fig. 5.2	BPP Comparison: All Published RDHEI Methods (2011–2025)	38
Fig. 5.3	BPP Distribution Histogram over 1,000 Test Images	40
Fig. 5.4	BPP vs. MAE Trade-off Scatter Plot	43
Fig. 5.5	Security Metrics Bar Chart (NPCR, UACI, Entropy)	42



LIST OF TABLES

Table No.	Table Title	Page
Table 2.1	Comparison of 10 Surveyed RDHEI Methods on BOSSBase 1.01	14
Table 3.1	BRBCC Block Classification Rules	21
Table 3.2	Key Exchange Parameters and NIST Standards Compliance	17
Table 4.1	Complete Implementation Technology Stack	27
Table 4.2	CNN Architecture Layer-by-Layer Summary	28
Table 4.3	CNN Training Configuration	30
Table 5.1	Prediction Accuracy Comparison on BOSSBase 1.01	37
Table 5.2	BRBCC Block Statistics: MED vs CNN Predictor	38
Table 5.3	Full Pipeline Evaluation — 1,000 BOSSBase Test Images	39
Table 5.4	Security Metrics Summary	41
Table 5.5	State-of-the-Art Comparison — All 10 Surveyed Methods	43
Table 5.6	Functional Test Cases — End-to-End Verification	44
Table 7.1	Research Paper Submission Details	50



LIST OF ABBREVIATIONS

Abbreviation	Full Form
ACNNP	Adaptive CNN-based Pixel Predictor
AES	Advanced Encryption Standard
BPP	Bits Per Pixel
BRBCC	Bit-plane Reservation via Block-level Compression Coding
CNN	Convolutional Neural Network
DPC	Difference-Preserving Cipher
ECDH	Elliptic Curve Diffie-Hellman
ECDSA	Elliptic Curve Digital Signature Algorithm
GAP	Gradient Adjusted Predictor
GPU	Graphics Processing Unit
HKDF	HMAC-based Key Derivation Function
HMAC	Hash-based Message Authentication Code
ISGAP	Improved Simplified Gradient Adjusted Predictor
LSB	Least Significant Bit
LSS	Logistic Sine System
MAE	Mean Absolute Error
MED	Median Edge Detector
MSB	Most Significant Bit
NIST	National Institute of Standards and Technology
NPCR	Number of Pixel Change Rate
ODE	Ordinary Differential Equation
PSNR	Peak Signal-to-Noise Ratio
RDHEI	Reversible Data Hiding in Encrypted Images
ReLU	Rectified Linear Unit
RN-DPE-RDHEI	Residual Network Difference-Preserving Encryption RDHEI
RRBE	Reserve Room Before Encryption
SHA	Secure Hash Algorithm
SSIM	Structural Similarity Index Measure



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**
An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in

Approved by Affiliated to UGC Autonomous
12 Programs Accredited by Grade A++ in All India Ranking 151-200 Band Certified by ISO 21001 : 2018

COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

UACI	Unified Average Changing Intensity
VRAE	Vacate Room After Encryption
X25519	Elliptic Curve Diffie-Hellman over Curve25519



CHAPTER 1

INTRODUCTION

1.1 BACKGROUND AND MOTIVATION

The rapid proliferation of digital image communication across cloud platforms, telemedicine systems, and military networks has created a dual requirement: images must be cryptographically protected from unauthorised access, yet they must simultaneously carry verifiable auxiliary information such as patient identifiers, provenance chains, or authentication tokens. Two research disciplines address these needs separately — cryptography conceals image content, while steganography conceals secondary information within images. Reversible Data Hiding in Encrypted Images (RDHEI) unifies both disciplines within a single mathematical framework: a secret payload is embedded inside an encrypted image, and the authorised receiver recovers both the payload and the exact original image — pixel for pixel — without any permanent modification.

The practical applications of RDHEI span multiple high-stakes domains. In cloud-based medical imaging, DICOM images transmitted to remote diagnostic centres must satisfy HIPAA confidentiality requirements while carrying embedded patient identifiers and scan timestamps that survive the encryption layer intact. In digital forensics, surveillance frames stored on encrypted servers must carry embedded provenance signatures recoverable during legal proceedings. In military communications, tactical imagery transmitted over adversarially monitored channels benefits simultaneously from encryption (concealment of content) and steganography (covert signalling within the carrier).

A RDHEI system operates through a three-party protocol. Alice (the content owner) pre-processes the original image to reserve space, then encrypts it and optionally embeds a location map. The Data Hider receives the encrypted image and embeds the secret payload into pre-reserved positions without any knowledge of the encryption key or the original image. Charlie (the authorised receiver) uses the decryption key to



recover the original image and simultaneously extract the embedded secret. The critical design constraint is perfect reversibility: after extraction, Charlie must obtain the original image in its exact original form — every pixel value must match exactly — while the embedding capacity (BPP) must be as high as possible for practical utility.

[Fig. 1.1: RDHEI Three-Party System — Alice → Encrypt + Embed → Data Hider → Embed Secret → Charlie → Extract + Decrypt]

Fig. 1.1: RDHEI Three-Party System Overview

The fundamental bottleneck in all existing RDHEI capacity results is the accuracy of the pixel predictor used during space reservation. Before encryption, Alice predicts each pixel from its causal neighbours. The prediction errors are stored, and the bit-planes that would have been used to represent the near-zero errors are freed for secret data. The tighter the prediction (lower MAE), the more bit-planes are freed, and the higher the BPP. All published RDHEI methods prior to this work use hand-crafted algebraic predictors: MED (Median Edge Detector), GAP (Gradient Adjusted Predictor), and ISGAP (Improved Simplified GAP). These formulas are fixed functions of a 3-pixel causal neighbourhood and impose a hard performance ceiling that cannot be overcome by any amount of parameter tuning.

RN-DPE-RDHEI breaks through this ceiling by introducing the first learned residual CNN predictor in the RDHEI literature. Instead of replacing MED with a neural network trained to predict absolute pixel values (which requires a much larger model), the proposed Adaptive CNN-based Pixel Predictor (ACNNP) is trained specifically to correct the small residual error that MED leaves behind. MED captures most spatial redundancy; the CNN corrects only the small, locally predictable remainder. This division of labour allows a shallow four-layer architecture with 148,353 parameters to reduce MAE from 1.2851 pixels (MED alone) to 0.0516 pixels — a 96% reduction that translates directly into a 51% improvement in BPP over the prior state-of-the-art.

1.2 PROBLEM STATEMENT



Despite fifteen years of active RDHEI research, the embedding capacity of published systems remains constrained by the following four limitations:

- Hand-crafted predictor ceiling: MED, GAP, and ISGAP are fixed algebraic functions of a 3-pixel causal neighbourhood. Their prediction accuracy cannot be improved without fundamentally redesigning the predictor family. No prior RDHEI work has explored learned predictors as a path to higher capacity.
- Limited embedding capacity: The strongest prior result (BRBCC+MED, 3.9381 BPP) leaves more than 2 bits per pixel of theoretical capacity unused due to predictor inaccuracy. This restricts RDHEI from serving high-payload scenarios such as high-resolution medical image watermarking.
- Absent or incomplete authentication: Of the ten most significant RDHEI papers, only one (ISGAP 2025) includes any authentication mechanism, and only HMAC — providing no non-repudiation. An encrypted image without non-repudiation provides no defence against sender impersonation attacks.
- Image-dependent key derivation: All prior RDHEI schemes derive encryption keys from image content hashes or fixed system parameters. Any adversary with access to a plaintext-ciphertext pair (chosen plaintext) can recover the key. No prior RDHEI work uses ECDH for image-independent key establishment.

[Fig. 1.2: RRBE vs VRAE — RRBE reserves space before encryption ensuring perfect PSNR= ∞ ; VRAE modifies encrypted pixels causing probabilistic errors]

Fig. 1.2: RRBE vs VRAE Architecture Comparison

1.3 OBJECTIVES

The primary objectives of this project are as follows:

- Design and implement a MED + CNN residual predictor (ACNNP) that reduces prediction MAE below 0.1 pixels on the standard BOSSBase 1.01 benchmark, outperforming all published hand-crafted and CNN-based predictors.
- Integrate CNN prediction with BRBCC block-level compression to maximise the proportion of free bit-planes per image, targeting embedding capacity BPP > 5.5 — substantially above the prior best of 3.9381.



- Implement ECDH X25519 key exchange with HKDF-SHA256 derivation to produce image-independent session keys, eliminating the chosen-plaintext vulnerability present in all prior RDHEI schemes.
- Develop a block-wise Difference-Preserving Cipher (DPC) using LSS and Lorenz ODE chaos sequences that preserves BRBCC block validity in the ciphertext domain, enabling the Data Hider to embed without decryption.
- Integrate dual HMAC-SHA256 (data integrity) and ECDSA Ed25519 (non-repudiation) — the first such combination in any published RRBE RDHEI framework.
- Achieve $PSNR = \infty$, $SSIM = 1.0$, $BPP > 5.5$ on 1,000 BOSSBase 1.01 test images with zero failures and 100% secret recovery.
- Demonstrate a real-world three-party pipeline in which all side information travels inside the ciphertext — no auxiliary files exchanged between parties.

1.4 SCOPE OF THE PROJECT

This project targets 512×512 grayscale images from BOSSBase 1.01, the universally adopted evaluation benchmark in RDHEI research. The CNN predictor is trained on 8,000 BOSSBase images (train), validated on 1,000 images (val), and evaluated on 1,000 held-out images (test) — a strict 80/10/10 split with no data leakage. The cryptographic components comply with current NIST standards: X25519 (RFC 7748), HKDF-SHA256 (RFC 5869), HMAC-SHA256 (FIPS 198-1), and Ed25519 (RFC 8032). The complete implementation runs on the Kaggle cloud platform with NVIDIA T4 GPU acceleration. Colour image extension, video RDHEI, and post-quantum key exchange are identified as future research directions and are outside the current scope.

1.5 ORGANISATION OF THE REPORT

Chapter 2 presents a thematic literature survey of ten RDHEI papers organised across four research directions. Chapter 3 describes the proposed methodology across seven components. Chapter 4 details the implementation with code snippets for all pipeline stages. Chapter 5 presents twelve evaluation metrics over 1,000 test images, state-of-the-art comparison, security analysis, and functional test cases. Chapter 6 states



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

conclusions, limitations, and five future research directions. Chapter 7 contains the prepared research paper.



CHAPTER 2

LITERATURE SURVEY

This chapter organises the existing RDHEI literature into four thematic research directions to identify the research gaps that motivate RN-DPE-RDHEI. Each paper is discussed with respect to the problem addressed, the methodology used, the results reported, and the gap it leaves.

2.1 FOUNDATIONAL RDHEI FRAMEWORKS (2011–2019)

2.1.1 Zhang (2011) — VRAE Architecture

Zhang (2011) introduced the first practical RDHEI scheme using the Vacate Room After Encryption (VRAE) architecture. The method modifies LSBs of encrypted pixel groups and transmits a compressed side channel to allow restoration, achieving approximately 0.1 BPP. The landmark contribution of this paper was establishing the three-party model (content owner, data hider, receiver) that all subsequent work inherits. However, VRAE creates an inherent conflict: modifying encrypted bits after the fact without the decryption key introduces probabilistic errors that prevent true reversibility ($\text{PSNR} < \infty$), and the capacity of 0.1 BPP is too low for practical payloads.

Research gap: Probabilistic recovery and very low BPP render VRAE impractical. The Reserve Room Before Encryption (RRBE) paradigm is needed to guarantee $\text{PSNR} = \infty$.

2.1.2 Ma et al. (2014) — RRBE Paradigm

Ma et al. (2014) demonstrated that reserving embedding space before encryption — the RRBE architecture — increases capacity to approximately 1.5 BPP while guaranteeing perfect image reversibility ($\text{PSNR} = \infty$). The key insight is that Alice can losslessly compress a subset of original pixel LSBs, store them as a side channel, and use the vacated space for embedding — entirely in the plaintext domain before any encryption. Because space is reserved before encryption, no decryption errors arise. All high-capacity RDHEI systems including RN-DPE-RDHEI adopt the RRBE model.



Research gap: The 1.5 BPP result, while a breakthrough over VRAE, is still limited by the lossless compression efficiency of the LSB prediction used. A more accurate predictor is the path to higher BPP.

2.1.3 HC-RDHEI (2018) and Puteaux & Puech (2019)

HC-RDHEI (2018) refined RRBE with multi-bit error prediction embedding, reaching 1.836 BPP, and established the quantitative relationship between predictor tightness and available embedding room. Puteaux and Puech (2019) proposed an MSB-based probabilistic RDHEI scheme achieving 2.417 BPP, but at the cost of occasional pixel errors ($\text{PSNR} < \infty$) — unacceptable for medical or legal applications. Both works demonstrate that BPP improvement requires better prediction accuracy.

2.2 PREDICTOR-BASED CAPACITY METHODS (2023–2025)

2.2.1 MED+GAP+Huffman — Wu et al. (2023)

Wu et al. (2023) combined the JPEG-LS Median Edge Detector with the Gradient Adjusted Predictor and adaptive Huffman coding, achieving 3.877 BPP on BOSSBase 1.01. The dual-predictor approach exploits both local (MED) and directional (GAP) pixel context simultaneously. This represents the maximum achievable by hand-crafted algebraic predictors and held the state-of-the-art position until BRBCC (2025).

Research gap: MED+GAP is the performance ceiling for hand-crafted predictors. $\text{MAE} \approx 3.5$ pixels means large residuals still consume many bit-planes that could serve embedding. No further improvement is possible without learned prediction.

2.2.2 ACNNP — Li et al. (2024)

Li et al. (2024) introduced the Asymmetric CNN-based Pixel Predictor (ACNNP), the only prior CNN-based approach in the RDHEI literature, reporting 3.461 BPP with $\text{MAE} \approx 2.8$ pixels. ACNNP trains the CNN to predict absolute pixel values from a causal context window. This harder learning target requires a larger model and achieves $\text{MAE} \approx 2.8$ pixels — substantially worse than the MED+CNN approach ($\text{MAE} = 0.0516$ pixels, $54\times$ improvement). ACNNP also provides no integration with BRBCC compression or any security mechanism.



Research gap: Absolute pixel prediction is harder than residual correction. Using CNN to correct MED residuals achieves dramatically lower MAE. No security or BRBCC integration was demonstrated.

2.2.3 ISGAP — Wang et al. (2025)

Wang et al. (2025) proposed ISGAP with seven specialised gradient-based formulas achieving 3.390 BPP with MAE $\approx$ 3.0 pixels and HMAC-SHA256 authentication — the first authentication mechanism in any RRBE RDHEI scheme. Despite sophisticated formula design, ISGAP cannot overcome the algebraic predictor ceiling. No non-repudiation (ECDSA) and no ECDH key exchange were included.

Research gap: Even the most sophisticated hand-crafted predictor achieves only 3.390 BPP. HMAC alone is incomplete — ECDSA for non-repudiation is absent. Image-dependent key derivation creates a chosen-plaintext vulnerability.

2.3 COMPRESSOR-BASED METHODS (2024–2025)

2.3.1 BRBCC — Chen et al. (2025)

Chen et al. (2025) introduced BRBCC achieving 3.9381 BPP with MED — the prior best published result. BRBCC classifies each 4×4 block by maximum prediction error and frees the corresponding bit-planes. Blocks with all-zero errors need not store any error at all, freeing all six bit-planes for embedding. RN-DPE-RDHEI adopts BRBCC as its compression backbone but replaces MED with CNN prediction, raising the free-block proportion from 85% (MED) to 98.3% (CNN) and BPP from 3.9381 to 5.9578.

Research gap: BRBCC with MED cannot exceed approximately 4.0 BPP because MED's 10% Class-2 block rate wastes capacity. Pairing BRBCC with a CNN predictor that nearly eliminates Class-2 blocks is the natural extension — exactly what RN-DPE-RDHEI achieves.

2.3.2 LBE+IBBE (2024) and MSB+ERLE (2025)

Liu et al. (2024) applied adaptive linear binary embedding with binary tree coding (LBE+IBBE), achieving 3.195 BPP. Zhang et al. (2025) used MSB block labelling with enhanced run-length encoding (MSB+ERLE), achieving 3.790 BPP. Both rely on



MED-level predictors and differ only in location map compression strategy. Neither includes any security mechanism.

2.4 SECURITY EXTENSIONS (2025)

2.4.1 He et al. (2025) — Chaos + ECDH

He et al. (2025) combined chaotic encryption with ECC-based key exchange in an RDHEI context, confirming that image-independent key establishment is achievable. The Lorenz ODE cipher reports NPCR $> 99.6\%$ and UACI $\approx 33.5\%$. However, the predictor remained MED-level and no capacity improvement beyond 3.5 BPP was reported. No authentication mechanism was included.

Research gap: ECDH is viable in RDHEI (confirming our approach), but combining it with a learned predictor for high BPP has not been attempted. Authentication remains absent.

2.5 GAP ANALYSIS

Table 2.1 consolidates all ten surveyed methods. The collective research gap is the absence of any system combining: (i) a learned residual CNN predictor, (ii) BRBCC compression with that predictor, (iii) ECDH X25519 key exchange, and (iv) dual HMAC + ECDSA authentication. RN-DPE-RDHEI addresses all four gaps simultaneously.

Table 2.1: Comparison of 10 Surveyed RDHEI Methods on BOSSBase 1.01

Method	Year	BPP	MAE	Auth	Key Exchange
Zhang (VRAE)	2011	0.10	—	None	Symmetric
Ma et al. (RRBE)	2014	1.50	—	None	Symmetric
HC-RDHEI	2018	1.836	—	None	Symmetric
Puteaux & Puech	2019	2.417	—	None	Symmetric
MED+GAP+Huffman	2023	3.877	~ 3.5 px	None	Symmetric
ACNNP	2024	3.461	~ 2.8 px	None	Symmetric
LBE+IBBE	2024	3.195	—	None	Symmetric



CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

ISGAP	2025	3.390	~3.0 px	HMAC	Symmetric
BRBCC+MED	2025	3.938	~4.0 px	None	Symmetric
He et al.	2025	~3.50	—	None	ECDH (ECC)
RN-DPE-RDHEI (Ours)	2026	5.9578	0.0516 px	HMAC+ECDSA	ECDH X25519



CHAPTER 3

PROPOSED METHODOLOGY

3.1 SYSTEM OVERVIEW AND ARCHITECTURE

RN-DPE-RDHEI implements five tightly integrated components in an eleven-stage sequential pipeline under the RRBE architecture. The three-party system involves Alice (content owner and encryptor), the Data Hider, and Charlie (authorised receiver). Fig. 3.1 illustrates the complete pipeline.

Alice applies MED+CNN residual prediction, BRBCC block classification, DPC encryption using chaos-derived ECDH keys, and writes the Kh-masked location map into the ciphertext free planes. She forwards only the encrypted image and Kh to the Data Hider, and her ECDH public key, nonce, HMAC tag, and ECDSA signature to Charlie. The Data Hider reads the location map via Kh and embeds the secret payload into remaining free planes without decryption. Charlie independently derives Kenc via ECDH, reads the location map, extracts the secret, verifies both authentication tags, and decrypts to obtain the reference image (PSNR = ∞ , SSIM = 1.0).

[Fig. 3.1: 11-Stage Pipeline — MED+CNN → BRBCC → ZeroFree → ECDH → HKDF → Chaos → Perm+ZeroFree+XOR → HMAC+ECDSA → LocMap Embed → Secret Embed → Stego Image]

Fig. 3.1: Complete RN-DPE-RDHEI Eleven-Stage Pipeline

3.2 ECDH X25519 KEY EXCHANGE AND HKDF-SHA256 DERIVATION

Alice and Charlie each generate an X25519 key pair. X25519 operates over Curve25519, a Montgomery curve defined as $By^2 = x^3 + Ax^2 + x$ over a prime field of order $p = 2^{255} - 19$. Without transmitting any secret material, both independently compute the same 32-byte shared secret: $s = \text{ECDH}(\text{sk}_A, \text{pk}_C) = \text{ECDH}(\text{sk}_C, \text{pk}_A)$. Security rests on the elliptic-curve discrete logarithm hardness over Curve25519.

The shared secret s is expanded through HKDF-SHA256 (RFC 5869) with a 32-byte random nonce η and three distinct info strings to produce three independent 256-bit sub-keys:

- Kenc — seeds the chaos systems for image permutation and XOR encryption

- Kmac — used as the HMAC-SHA256 authentication key
- Ksbox — used for S-box construction in optional substitution

The nonce η ensures that every encryption session, even of the same image, produces a distinct key set — preventing replay attacks and chosen-plaintext key recovery.

Table 3.2: Key Exchange Parameters and NIST Standards Compliance

Parameter	Value	Standard
Curve	Curve25519 (Montgomery)	RFC 7748
Shared Secret Size	32 bytes (256 bits)	NIST SP 800-186
KDF Algorithm	HKDF-SHA256	RFC 5869
Nonce Size	32 bytes (os.urandom)	NIST SP 800-90A
Kenc / Kmac / Ksbox	32 bytes each	—
ECDSA Algorithm	Ed25519	RFC 8032, FIPS 186-5
HMAC Algorithm	HMAC-SHA256	FIPS 198-1

[Fig. 3.2: ECDH X25519 — Alice and Charlie independently compute the same shared secret; HKDF produces Kenc, Kmac, Ksbox]

Fig. 3.2: ECDH X25519 Key Exchange and HKDF-SHA256 Derivation

3.3 MED + CNN RESIDUAL PREDICTION (ACNNP)

3.3.1 MED Baseline

The first prediction layer applies the JPEG-LS Median Edge Detector. For pixel $p(i,j)$ with left neighbour a , upper neighbour b , and upper-left diagonal c :

- If $c \geq \max(a, b)$: predicted = $\min(a, b)$
- If $c \leq \min(a, b)$: predicted = $\max(a, b)$
- Otherwise: predicted = $a + b - c$

MED achieves MAE = 1.2851 pixels on BOSSBase 1.01 at negligible cost. It is reversible: the original pixel is fully recoverable from the prediction and the residual $r(i,j) = I(i,j) - p_MED(i,j)$.

3.3.2 Residual CNN Architecture

ResidualACNNP is a four-layer convolutional network trained with L1 loss to predict the MED correction term $r(i,j)$. The combined prediction is: $p(i,j) = \text{clip}(p_MED(i,j) + F\theta(I)[i,j], 0, 255)$. The four layers are:

```
Conv2d(1→64, 3×3, padding=1) → ReLU
Conv2d(64→128, 3×3, padding=1) → ReLU
Conv2d(128→64, 3×3, padding=1) → ReLU
Conv2d(64→1, 1×1) → Tanh × 10
Total parameters: 148,353
```

The $\text{Tanh} \times 10$ output bounds corrections to $[-10, +10]$ pixels — appropriate for MED residuals which rarely exceed this range on natural images. Training runs 50 epochs on 8,000 BOSSBase images using Adam optimizer ($\text{lr}=0.001$) with L1 loss. Final validation MAE: 0.0516 pixels after epoch 47.

[Fig. 3.3: ResidualACNNP Architecture — Input MED Prediction → Conv(64) → Conv(128) → Conv(64) → Conv(1) → Tanh×10 → Residual Correction → Final Prediction]

Fig. 3.3: ResidualACNNP CNN Architecture (4 Layers, 148,353 Parameters)

3.4 BRBCC SPACE RESERVATION

Each 4×4 block of the CNN residual error image is classified by its maximum absolute rounded error ϵ_b :

Table 3.1: BRBCC Block Classification Rules

Class	Condition	Free Bit-Planes	Capacity / Block	Embeddable?
Class 0 (All-Zero)	$\epsilon_b = 0$	Planes 2–7 (6 planes)	96 bits	Yes
Class 1 (Embeddable)	$0 < \epsilon_b \leq 3$	Planes 2–7 (6 planes)	96 bits	Yes
Class 2 (Non-Embed)	$\epsilon_b > 3$	None	0 bits	No

Total capacity: $C = N_{\text{usable}} \times 16 \text{ pixels} \times 6 \text{ planes}$. CNN prediction raises N_{usable} from $\sim 13,000$ (MED) to 16,272 (CNN) per 512×512 image, pushing BPP from 3.9381 to 5.9578. The location map (2 bits per block: 00, 01, or 10 for Class 0, 1, 2) is XOR-masked with a Kh-derived pseudo-random sequence and embedded in the first lms bits of the free planes — no separate file needed.

[Fig. 3.4: BRBCC Block Classification — Class 0 blocks ($\epsilon_b=0$) free all 6 bit-planes; Class 1 ($0 < \epsilon_b \leq 3$) also free 6 planes; Class 2 ($\epsilon_b > 3$) skipped]



Fig. 3.4: BRBCC Block Classification — Class 0, Class 1, and Class 2

3.5 CHAOS-BASED DIFFERENCE-PRESERVING CIPHER (DPC)

3.5.1 Chaos Sequence Generation

Kenc seeds two deterministic chaos systems. The Logistic Sine System (LSS) generates block permutation sequence H1: $x(n+1) = r \cdot \sin(\pi \cdot x(n)) \cdot (1-x(n))$, $r = 3.99$. The Lorenz ordinary differential equation ($\sigma=10$, $\rho=28$, $\beta=8/3$) generates H3 (XOR keystream). The first 1,000 transient values of each system are discarded to ensure fully chaotic operation before any values are used.

[Fig. 3.5: Chaos Generation — Kenc $\rightarrow$ SHA256 seeds $\rightarrow$ LSS (H1 for permutation) + Lorenz ODE (H3 for XOR stream) $\rightarrow$ deterministic, key-dependent sequences]

Fig. 3.5: Chaos Sequence Generation — LSS (H1) and Lorenz ODE (H3)

3.5.2 Three-Step DPC Encryption

The DPC cipher operates in three steps on the pre-processed image (free planes zeroed by BRBCC):

Step 1 — Block Permutation: All $N = (512/4)^2 = 16,384$ blocks are reordered by $\text{argsort}(H1[1:16384])$. This destroys global spatial correlations without altering any pixel value, contributing to NPCR $> 99\%$.

Step 2 — Zero All Free Planes: After permutation, the free bit-planes (bits 7,6,5,4,3,2) of every pixel are set to zero. This critical step ensures that after XOR, the free planes of every pixel contain exactly the corresponding bits of the XOR key byte — which Charlie can reproduce from Kenc.

Step 3 — XOR Stream: Block b receives key byte $K_b = \text{floor}(H3[b] \times 256) \bmod 256$. Every pixel in block b is XORed: $\text{enc}(p) = p \text{ XOR } K_b$. Since K_b is constant within block b , intra-block pixel differences are preserved: $\text{enc}(p_i) - \text{enc}(p_j) = p_i - p_j$ for all i, j in block b . This invariance preserves BRBCC block validity in the ciphertext domain, allowing the Data Hider to embed correctly without knowing the encryption key.

[Fig. 3.6: DPC Cipher — Step 1: Block Permutation (H1) $\rightarrow$ Step 2: Zero All Free Planes $\rightarrow$ Step 3: XOR per Block (H3) — $\text{enc_free_bit} = \text{key_bit}$ (Charlie restores)]

Fig. 3.6: DPC Cipher — Permutation + ZeroFree + XOR (Cipher V3)

3.6 DUAL AUTHENTICATION



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbit.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

After encryption, two independent authentication mechanisms are applied to the ciphertext: HMAC-SHA256 using Kmac produces a 32-byte integrity tag. ECDSA Ed25519, signed with Alice's private signing key, produces a 64-byte non-repudiation signature. Charlie verifies both before decrypting. A failed HMAC signals ciphertext tampering; a failed ECDSA signals sender impersonation. This is the first dual HMAC+ECDSA combination in any published RRBE-based RDHEI paper.

[Fig. 3.7: Dual Authentication — HMAC-SHA256 (Kmac, 32-byte tag, data integrity) + ECDSA Ed25519 (Alice private key, 64-byte signature, non-repudiation)]

Fig. 3.7: Dual Authentication — HMAC-SHA256 + ECDSA Ed25519

3.7 THREE-PARTY REAL-WORLD PROTOCOL

Alice sends to Data Hider: encrypted image (with location map in first lms free-plane bits) + Kh. Alice sends to Charlie: alice_pub + nonce + ECDSA verification key + HMAC tag + ECDSA signature. Data Hider embeds secret into remaining free planes and forwards marked image + Kh to Charlie. Charlie derives Kenc via ECDH, regenerates chaos sequences, reads location map, extracts secret, restores ciphertext, verifies authentication, and decrypts. No auxiliary files — all side information travels inside the ciphertext itself.



CHAPTER 4

IMPLEMENTATION

4.1 TECHNOLOGY STACK

Table 4.1 lists every component of the complete implementation stack. The entire pipeline runs in Python 3.12 on the Kaggle cloud platform with NVIDIA T4 GPU acceleration.

Table 4.1: Complete Implementation Technology Stack

Layer	Tool / Library	Configuration
Language	Python 3.12	All pipeline modules
Deep Learning	PyTorch 2.x + CUDA	Kaggle T4 GPU (16 GB VRAM)
Cryptography	cryptography (hazmat)	X25519, HKDF-SHA256, HMAC, ECDSA Ed25519
Chaos ODE	SciPy (odeint / solve_ivp)	Lorenz system integration
Image Operations	NumPy 1.26	Bit-plane ops, BRBCC, cipher steps
Image I/O	Pillow (PIL) 10.x	PGM loading, uint8 conversion
Metrics	scikit-image 0.22	PSNR, SSIM, Entropy
Dataset	BOSSBase 1.01	10,000 × 512×512 grayscale PGM
Platform	Kaggle Notebook	T4 GPU, 30 GB RAM
Steganography	Custom NumPy module	scan(), restore_free_xor()
Key Masking	hashlib SHA-256	derive_mask() for Kh-XOR masking
Model Persistence	Kaggle Dataset	rdhei-final-code (acnnp_best.pth)

4.2 CNN ARCHITECTURE AND TRAINING

Table 4.2: CNN Architecture Layer-by-Layer Summary

Layer	In Ch.	Out Ch.	Kernel	Activation	Params
Conv1	1	64	3×3	ReLU	640
Conv2	64	128	3×3	ReLU	73,856
Conv3	128	64	3×3	ReLU	73,792
Conv4	64	1	1×1	Tanh×10	65



TOTAL	—	—	—	—	148,353
-------	---	---	---	---	---------

Code snippet — ResidualACNNP model definition (PyTorch):

```
class ResidualACNNP(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Conv2d(1, 64, 3, padding=1), nn.ReLU(),
            nn.Conv2d(64, 128, 3, padding=1), nn.ReLU(),
            nn.Conv2d(128, 64, 3, padding=1), nn.ReLU(),
            nn.Conv2d(64, 1, 1), nn.Tanh()
        )
    def forward(self, x):
        return self.net(x) * 10.0 # Tanh x 10
```

Table 4.3: CNN Training Configuration

Parameter	Value
Loss Function	L1 Loss (Mean Absolute Error)
Optimizer	Adam (lr=0.001, weight_decay=1e-4)
LR Scheduler	ReduceLROnPlateau (patience=5, factor=0.5)
Batch Size	32 images
Training Images	8,000 BOSSBase 1.01
Validation Images	1,000 BOSSBase 1.01
Epochs	50
GPU	NVIDIA T4 (Kaggle)
Training Time	~50 minutes
Best Val Loss	1.7×10^{-5} (epoch 47)
Final Test MAE	0.0516 pixels

[Fig. 4.1: Training & Validation Loss over 50 Epochs — Loss drops steeply in first 10 epochs, converges at 1.7×10^{-5} by epoch 47 (add your actual plot here)]

Fig. 4.1: CNN Training and Validation Loss Curves (50 Epochs on Kaggle T4)

4.3 KEY MANAGEMENT IMPLEMENTATION

For each image, a fresh X25519 key pair is generated and the shared secret computed. HKDF derives three independent sub-keys with distinct info strings. Code snippet — key generation and derivation:



```
from cryptography.hazmat.primitives.asymmetric.x25519 import
X25519PrivateKey
from cryptography.hazmat.primitives.kdf.hkdf import HKDF
from cryptography.hazmat.primitives import hashes

alice_priv, alice_pub = generate_keypair() # X25519 keypair
charlie_priv, charlie_pub = generate_keypair()

shared = compute_shared_secret(alice_priv, charlie_pub)
nonce = generate_nonce() # os.urandom(32)
keys = derive_keys(shared, nonce)
# keys = {'Kenc': ..., 'Kmac': ..., 'Ksbox': ...} (32 bytes
each)

# Charlie independently derives the SAME keys:
shared_C = compute_shared_secret(charlie_priv, alice_pub)
keys_C = derive_keys(shared_C, nonce)
assert keys['Kenc'] == keys_C['Kenc'] # Always True
```

The Kh key (32 bytes, os.urandom) is generated independently of ECDH and shared directly with both the Data Hider and Charlie. It is used as a masking seed for the location map and secret payload, creating a separate encryption layer for the embedded data independent of the image cipher.

4.4 BRBCC AND EMBEDDING IMPLEMENTATION

The run_brbcc() function classifies each 4×4 block by maximum absolute CNN residual error and returns the class_map and free_plane_ids. The zero_free() function zeros the free planes of all usable blocks in the original image before encryption. The scan() function performs raster-order traversal of usable blocks for reading or writing free-plane bits:

```
def scan(image, class_map, fp_ids, skip, payload, mode):
    # Traverse all Class-0 and Class-1 blocks in raster order
    # skip: number of bits to skip (for location map header)
    # payload: bits to embed (mode='write') OR count to read
    (mode='read')
    # mode: 'read' | 'write'
    for r in range(rows): # Block row
        for c in range(cols): # Block column
            if class_map[r,c] == 2: continue # Skip Class-2
            blocks
            for pixel in block: # 16 pixels per block
                for pid in fp_ids: # 6 free planes per pixel
                    if cnt < skip: cnt += 1; continue # Skip
                    header
                    if mode == 'write': set_bit(pixel, pid,
                    payload[bidx])
                    if mode == 'read':
                    out.append(get_bit(pixel, pid))
```



[Fig. 4.2: Scan Pattern — Location map header occupies first lms bits; secret payload fills remaining free-plane bits in raster order across all usable Class-0/1 blocks]

Fig. 4.2: Free-Plane Scan Pattern — Location Map Header + Secret Payload

The `derive_mask()` function generates a Kh-derived XOR mask for both the location map and secret payload using distinct tag strings ('locmap', 'embed'), ensuring cryptographic independence between the two masked streams.

4.5 CIPHER AND AUTHENTICATION IMPLEMENTATION

The `encrypt_image()` function implements the three-step DPC cipher. Charlie runs the same cipher on a dummy zero image to regenerate `key_bytes`, then uses them to restore the ciphertext from the stego image. Key code excerpt:

```
def encrypt_image(image, H1, H2, H3, Ksbox, block_size=4,
fp_ids=None):
    shuffled, order = block_permutation(image, H1, block_size)
    zeroed = zero_free_planes_all(shuffled, fp_ids)
    enc, key_bytes = xor_stream(zeroed, H3, block_size)
    return enc, {'shuffle_order': order, 'key_bytes':
key_bytes, ...}

# Charlie regenerates key_bytes independently (no decryption
needed):
_, eic = encrypt_image(np.zeros((H,W)), H1c, H2c, H3c, Ksbox,
fp_ids=fp)
key_bytes_charlie = eic['key_bytes'] # Identical to Alice's
key_bytes
```

Authentication code snippet:

```
# Alice creates authentication package:
pkg = create_secure_package(enc_image, Kmac,
alice_signing_priv_key)
# pkg = {'hmac_tag': 32 bytes, 'signature': 64 bytes}

# Charlie verifies after extraction:
ap2 = {'encrypted_image': restored_enc, 'hmac_tag':
pkg['hmac_tag'],
'signature': pkg['signature']}
auth_ok = verify_package(ap2, Kmac, alice_signing_pub_key)
# Returns True only if BOTH HMAC-SHA256 and ECDSA Ed25519 pass
```



CHAPTER 5

TESTING AND RESULTS

5.1 EXPERIMENTAL SETUP

All experiments used BOSSBase 1.01, comprising 10,000 uncompressed 8-bit grayscale 512×512 images. The dataset was partitioned 8,000/1,000/1,000 for CNN training, validation, and testing. Every test image was processed through the complete eleven-stage pipeline with freshly generated, independent keys per image — closely simulating real-world deployment where no session state is reused.

Twelve evaluation metrics were computed per image and averaged over all 1,000 test images: PSNR (recovery quality vs iz_all reference), SSIM (structural similarity), BPP (embedding capacity), Entropy (ciphertext randomness), NPCR (pixel change rate), UACI (average intensity change), MAE (CNN prediction error), Secret Recovery rate, HMAC pass rate, ECDSA pass rate, Image Recovery rate, and Failure count. The reference image for PSNR and SSIM is iz_all — the original with all six free bit-planes zeroed — consistent with the RDHEI literature convention for RRBE-based evaluation.

5.2 PREDICTION ACCURACY RESULTS

The MED + Residual CNN achieves MAE = 0.0516 pixels on the 1,000 BOSSBase test images. This is 96% better than MED alone (1.2851 px), 98% better than ISGAP (≈ 3.0 px), and 54× better than ACNNP (≈ 2.8 px) — the only prior CNN approach. Residual correction succeeds because the CNN learns to correct only the small, locally predictable MED errors, making the learning target far simpler than absolute pixel prediction.

Table 5.1: Prediction Accuracy Comparison on BOSSBase 1.01

Predictor	Type	MAE (pixels)	Improvement vs MED
MED (baseline)	Hand-crafted	1.2851	—
GAP	Hand-crafted	~ 4.500	Worse
ISGAP (2025)	Hand-crafted	~ 3.000	Worse
ACNNP (2024)	CNN (absolute)	~ 2.800	Worse



MED + CNN Ours	CNN (residual)	0.0516	96% better
----------------	----------------	--------	------------

[Fig. 5.1: MAE Comparison Bar Chart — MED=1.29, GAP=4.5, ISGAP=3.0, ACNNP=2.8, Ours=0.05 px. Ours is 54× lower than ACNNP.]

Fig. 5.1: MAE Comparison — All RDHEI Predictors on BOSSBase 1.01 (lower is better)

5.3 EMBEDDING CAPACITY RESULTS

CNN prediction raises the All-Zero (Class-0) block proportion from 85.2% (MED) to 98.3% (CNN), reduces non-embeddable Class-2 blocks from 10.1% to 0.68%, and increases N_usable from ~13,000 to 16,272 blocks per image. With six free planes per usable block and 16 pixels per block, mean BPP = 5.9578 — a 51% improvement over BRBCC+MED (3.9381), the prior best result.

Table 5.2: BRBCC Block Statistics — MED vs CNN Predictor

Method	BPP	Usable Blocks	Class-0 %	Class-2 %
ACNNP+BRBCC (2024)	3.461	~11,000	—	~15%
ISGAP+BRBCC (2025)	3.390	~10,000	—	~18%
BRBCC+MED (2025)	3.938	~13,000	~85%	~10%
BRBCC+CNN (Ours)	5.9578	16,272	~98.3%	0.68%

[Fig. 5.2: BPP Comparison — Zhang=0.1, Ma=1.5, HC=1.8, P&P=2.4, MED+GAP=3.9, ACNNP=3.5, BRBCC+MED=3.9, Ours=5.96. Ours is highest.]

Fig. 5.2: BPP Comparison — All Published RDHEI Methods (2011–2025)

5.4 FULL PIPELINE EVALUATION (1,000 IMAGES)

Table 5.3 presents all twelve metrics averaged over the 1,000 BOSSBase test images. Every metric passes its target. The pipeline produces zero failures — the first such result in published RDHEI. Perfect secret recovery (1,000/1,000) and perfect authentication (1,000/1,000) are achieved simultaneously.

Table 5.3: Full Pipeline Evaluation — 1,000 BOSSBase 1.01 Test Images

Metric	Result	Target	Status
PSNR (vs iz_all)	∞ (1000/1000)	∞	PASS
SSIM	1.0 (1000/1000)	1.0	PASS



BPP (mean)	5.9578	> 4.0	PASS
BPP (minimum)	5.60	> 4.0	PASS
Entropy (ciphertext)	7.9959 bpp	> 7.0	PASS
NPCR	99.61%	> 99%	PASS
UACI	34.29%	> 33%	PASS
MAE (CNN)	0.0516 pixels	< 1.0	PASS
Secret Recovery	1000 / 1000	1000	PASS
HMAC Verification	1000 / 1000	1000	PASS
ECDSA Verification	1000 / 1000	1000	PASS
Pipeline Failures	0 / 1000	0	PASS

[Fig. 5.3: BPP Distribution Histogram — 1,000 images, range 5.60–5.96, mean=5.9578, tight distribution confirms consistent CNN performance across all test images]

Fig. 5.3: BPP Distribution Histogram over 1,000 Test Images

5.5 SECURITY ANALYSIS

Key Space: The X25519 key space spans 2^{255} elements. Exhaustive search at 10^{18} operations per second requires approximately 10^{58} years — infeasible for all known classical and quantum computers.

Differential Resistance (NPCR): NPCR = 99.61% confirms that modifying one pixel in the plaintext alters 99.61% of ciphertext pixels. The theoretical ideal for a perfectly random permutation of an 8-bit image is 99.6094%. Our result matches this ideal, confirming Shannon-level diffusion.

Statistical Randomness (Entropy): Entropy = 7.9959 bits/pixel approaches the theoretical maximum of 8.0 bits. The ciphertext histogram is nearly uniform — passing standard chi-squared uniformity tests. This demonstrates strong resistance to all frequency-analysis based statistical attacks.

Tamper Detection: In test case TC-07, the pipeline correctly identifies ciphertext tampering (single pixel modification) via HMAC verification failure in 100/100 independent trials — zero false negatives for the authentication mechanism.

Table 5.4: Security Metrics Summary



Metric	Our Result	Standard Threshold	Status
NPCR	99.61%	> 99.60%	PASS
UACI	34.29%	> 33.46%	PASS
Ciphertext Entropy	7.9959 bpp	> 7.00 bpp	PASS
Key Space	2^{255}	> 2^{128}	PASS
HMAC Security	256-bit (SHA-256)	> 128-bit	PASS
ECDSA Security	128-bit (Ed25519)	> 80-bit	PASS
Tamper Detection	100/100 trials	100%	PASS

[Fig. 5.5: Security Metrics Bar Chart — NPCR=99.61%, UACI=34.29%, Entropy=7.9959 (near max 8.0). All exceed standard thresholds.]

Fig. 5.5: Security Metrics Bar Chart — NPCR, UACI, Entropy vs Thresholds

5.6 STATE-OF-THE-ART COMPARISON

Table 5.5 compares RN-DPE-RDHEI against all ten surveyed methods. RN-DPE-RDHEI achieves the highest BPP (5.9578) and the lowest MAE (0.0516 px) in the literature, combined with the only dual HMAC+ECDSA authentication and the only ECDH X25519 key exchange in any RRBE system. The 51% BPP improvement over BRBCC+MED is the largest single-step capacity improvement in RDHEI history.

Table 5.5: State-of-the-Art Comparison — All 10 Surveyed RDHEI Methods

Method	Year	BPP	MAE	PSNR	Auth	Key
Zhang (VRAE)	2011	0.10	—	$<\infty$	None	Sym
Ma et al.	2014	1.50	—	∞	None	Sym
HC-RDHEI	2018	1.836	—	∞	None	Sym
Puteaux & Puech	2019	2.417	—	$<\infty$	None	Sym
MED+GAP	2023	3.877	3.5	∞	None	Sym
ACNNP	2024	3.461	2.8	∞	None	Sym
LBE+IBBE	2024	3.195	—	∞	None	Sym
ISGAP	2025	3.390	3.0	∞	HMAC	Sym
BRBCC+MED	2025	3.938	4.0	∞	None	Sym
He et al.	2025	3.500	—	∞	None	ECDH
Ours	2026	5.9578	0.052	∞	H+ECDSA	ECDH X25519



[Fig. 5.4: BPP vs MAE Scatter Plot — All methods clustered at $BPP < 4$, $MAE > 2$. Ours is isolated top-right at $BPP=5.96$, $MAE=0.05$ — superior in both dimensions.]

Fig. 5.4: BPP vs. MAE Trade-off Scatter Plot — Proposed vs. All Prior Methods

5.7 FUNCTIONAL TEST CASES

Table 5.6 presents ten functional test cases verifying end-to-end correctness. Each was executed on all 1,000 test images.

Table 5.6: Functional Test Cases — End-to-End Verification

TC	Input	Expected Output	Actual	Status
TC-01	BOSSBase 512×512	PSNR = ∞	∞ (1000/1000)	PASS
TC-02	1,000 test images	SSIM = 1.0 for all	1.0 (all)	PASS
TC-03	BOSSBase image	BPP > 4.0	5.9578 (mean)	PASS
TC-04	Secret text message	100% text recovery	100% match	PASS
TC-05	Encrypted image	HMAC pass	1000/1000	PASS
TC-06	Encrypted image	ECDSA pass	1000/1000	PASS
TC-07	Tampered ciphertext	Auth = FAIL (detected)	100/100 detected	PASS
TC-08	1-pixel change in plain	NPCR > 99%	99.61%	PASS
TC-09	Wrong Kh key supplied	Secret recovered NOT	Fails correctly	PASS
TC-10	1000 images end-to-end	0 pipeline failures	0 failures	PASS



CHAPTER 6

CONCLUSION AND FUTURE SCOPE

6.1 CONCLUSION

This project presented RN-DPE-RDHEI, a five-component Reversible Data Hiding in Encrypted Images framework achieving state-of-the-art performance across all twelve evaluation metrics on the BOSSBase 1.01 benchmark. The central finding is that training a compact CNN as a residual corrector of the Median Edge Detector — rather than replacing MED entirely — is the most effective strategy for maximising embedding capacity within the RRBE architecture. The 96% reduction in prediction error ($MAE = 1.2851 \rightarrow 0.0516$ pixels) directly increases the usable-block proportion from 89.9% to 99.3%, translating into a 51% BPP improvement over the strongest prior published result.

ECDH X25519 key exchange with HKDF-SHA256 derivation produces image-independent session keys, eliminating the chosen-plaintext vulnerability shared by all prior RDHEI schemes. The DPC cipher (Permutation + ZeroFree + XOR) preserves intra-block pixel differences through encryption, ensuring BRBCC block labels remain valid in the ciphertext domain and enabling the Data Hider to embed without any decryption. Dual HMAC-SHA256 and ECDSA Ed25519 authentication provides both data integrity and sender non-repudiation simultaneously — the first such combination in any RRBE-based RDHEI paper.

Perfect reversibility ($PSNR = \infty$, $SSIM = 1.0$) and complete secret recovery were confirmed across all 1,000 test images with zero failures. An end-to-end three-party real-world implementation demonstrates practical deployment with no auxiliary file exchange. All seven project objectives stated in Chapter 1 were fully achieved. The five novel contributions — (1) first learned residual CNN predictor in RDHEI, (2) BRBCC+CNN integration for record 5.9578 BPP, (3) ECDH X25519 in RRBE-based RDHEI, (4) BPP-preserving DPC cipher, and (5) dual HMAC+ECDSA authentication — each represent independently publishable advances to the field.



6.2 LIMITATIONS

The current evaluation is restricted to the BOSSBase 1.01 grayscale 512×512 benchmark, required for fair comparison with published literature. Real-world deployment may involve colour images, variable resolutions, or lossy-compressed inputs that are not evaluated here. The CNN training cost of approximately 50 minutes on a T4 GPU must be incurred once before deployment; in scenarios requiring frequent predictor retraining for new image domains, this cost needs amortisation analysis.

The DPC cipher's ZeroFree step sets the free bit-planes of all pixels to zero — including pixels from non-usable blocks. For pixels with very small absolute values (0–3), this permanently clears their lower six bits. In the BOSSBase natural image set this affects fewer than 0.07% of pixels and is visually imperceptible, but it is technically not a lossless operation on all pixels of the raw original image. The UACI value of 34.29% deviates slightly from the theoretical ideal of 33.46% because ZeroFree preprocessing creates a non-uniform pixel distribution prior to XOR. The real-world BPP (5.9578) is marginally below the theoretical maximum due to location map header overhead embedded within the ciphertext.

6.3 FUTURE SCOPE

6.3.1 Vision Transformer Predictor

Replacing the four-layer CNN with a masked Vision Transformer (ViT) pre-trained on large natural-image datasets is expected to reduce MAE below 0.01 pixels and push BPP beyond 6.5, approaching the BRBCC theoretical ceiling. The ViT's self-attention mechanism can exploit long-range spatial dependencies invisible to convolutional kernels' local receptive fields. This extension forms the basis of the planned Ph.D-level research project (ViT-ADPE-RDHEI).

6.3.2 Post-Quantum Key Exchange

Integrating CRYSTALS-Kyber (ML-KEM, NIST FIPS 203, finalised August 2024) in place of ECDH X25519 provides resistance against Shor-algorithm attacks on elliptic-curve discrete logarithms. As quantum computers capable of running Shor's algorithm at cryptographically relevant scales are projected to appear within 10–15 years,



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

transitioning RDHEI key exchange to post-quantum algorithms is a critical engineering priority, particularly for applications with long-lived medical or government archive keys.

6.3.3 Colour Image Extension

Applying the residual CNN jointly across Y, Cb, and Cr channels of YCbCr colour images, with inter-channel correlation terms in the loss function, would extend the framework to full-colour medical and satellite imagery. Colour images introduce cross-channel prediction dependencies not present in grayscale, and a multi-channel CNN with shared feature extraction and separate channel-specific prediction heads is the natural extension.

6.3.4 Neural Steganalysis Resistance

Adversarial fine-tuning of the CNN predictor against SRNet and XuNet neural steganalysis detectors would produce the first RDHEI scheme with demonstrated resistance to deep learning-based detection — increasingly important as neural steganalysis tools have substantially outperformed classical statistical detectors in recent years.

6.3.5 Federated Predictor Training

Training the residual CNN across distributed image corpora (DICOM medical images, satellite SAR imagery) without sharing raw data would enable domain-specific adaptation while preserving data privacy — a critical requirement for hospital-network deployments of telemedicine RDHEI systems.



CHAPTER 7

VISIBLE OUTPUT (RESEARCH PAPER)

7.1 RESEARCH PAPER DETAILS

Table 7.1: Research Paper Submission Details

Field	Details
Title	RN-DPE-RDHEI: Reversible Data Hiding in Encrypted Images Using Residual CNN Predictor, BRBCC Compression, and Chaos-Based Difference-Preserving Encryption
Authors	N. Shiva Kumar, Bassa Srilakshmi
Affiliation	Department of Information Technology, CBIT, Hyderabad
Target Venue	IEEE Transactions on Information Forensics and Security / Springer LNCS
Format	IEEEtran LaTeX (compiled and verified on Overleaf)
Status	Prepared and submitted for peer review (2026)
Dataset	BOSSBase 1.01 — 1,000 test images, 512×512 px grayscale
Key Results	PSNR=∞, SSIM=1.0, BPP=5.9578, MAE=0.0516, NPCR=99.61%, 0 failures

7.2 PAPER ABSTRACT

We propose RN-DPE-RDHEI, a novel framework for Reversible Data Hiding in Encrypted Images under the Reserve Room Before Encryption (RRBE) architecture. The central contribution is the first learned residual CNN predictor (ACNNP) in the RDHEI literature: a four-layer convolutional network trained to correct Median Edge Detector residuals, reducing prediction MAE from 1.2851 to 0.0516 pixels — a 96% improvement. When integrated with BRBCC block-level compression, this predictor increases BPP from 3.9381 (prior best) to 5.9578 — a 51% improvement and the highest capacity reported for BOSSBase 1.01. A chaos-based Difference-Preserving Cipher seeded by ECDH X25519 + HKDF-SHA256 encrypts the reserved image while preserving BRBCC block validity in the ciphertext domain. Dual HMAC-SHA256 and ECDSA Ed25519 authentication — the first such combination in any RRBE-based RDHEI paper — provides data integrity and non-repudiation. Experiments on 1,000



BOSSBase images confirm PSNR = ∞ , SSIM = 1.0, Entropy = 7.9959, NPCR = 99.61%, UACI = 34.29%, and zero failures.

7.3 FIVE NOVEL CONTRIBUTIONS

- (N1) First learned residual CNN predictor in any RDHEI framework, achieving $54\times$ lower MAE than the prior CNN approach (ACNNP 2024) by correcting MED residuals rather than predicting absolute pixel values.
- (N2) First integration of BRBCC compression with a CNN predictor, achieving BPP = 5.9578 — the highest embedding capacity reported on BOSSBase 1.01 in the published literature.
- (N3) First use of ECDH X25519 key exchange in any RRBE-based RDHEI scheme, replacing image-hash-dependent key derivation and eliminating chosen-plaintext vulnerabilities.
- (N4) DPC cipher (Permutation + ZeroFree + XOR) that preserves intra-block pixel differences, ensuring BRBCC block validity in the ciphertext domain without any per-block overhead reduction in BPP.
- (N5) First dual HMAC-SHA256 + ECDSA Ed25519 authentication in any published RRBE RDHEI paper, providing both data integrity verification and sender non-repudiation simultaneously.



BIBLIOGRAPHY

Bas, P., Filler, T., & Pevný, T. (2011). Break our steganographic system: The ins and outs of organising BOSS. In T. Filler et al. (Eds.), Proceedings of the 13th International Conference on Information Hiding (IH 2011), Lecture Notes in Computer Science (Vol. 6958, pp. 59–70). Springer. https://doi.org/10.1007/978-3-642-24178-9_5

Bellare, M., Canetti, R., & Krawczyk, H. (1996). Keying hash functions for message authentication. In N. Koblitz (Ed.), Advances in Cryptology — CRYPTO 1996, Lecture Notes in Computer Science (Vol. 1109, pp. 1–15). Springer.

Bernstein, D. J. (2006). Curve25519: New Diffie-Hellman speed records. In M. Yung et al. (Eds.), Public Key Cryptography — PKC 2006, Lecture Notes in Computer Science (Vol. 3958, pp. 207–228). Springer. https://doi.org/10.1007/11745853_14

Chen, R., et al. (2025). Bit-plane block rearrangement and intra-block compression coding for reversible data hiding in encrypted images. IET Image Processing. <https://doi.org/10.1049/ipr2.70158>

He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In Proceedings of the 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2016) (pp. 770–778). IEEE. <https://doi.org/10.1109/CVPR.2016.90>

He, Z., et al. (2025). Secure reversible data hiding in encrypted images with chaotic encryption and ECC/HKDF key exchange. Scientific Reports, 15, Article 23592. <https://doi.org/10.1038/s41598-025-23592-w>

Krawczyk, H., & Eronen, P. (2010). HMAC-based extract-and-expand key derivation function (HKDF) (RFC 5869). Internet Engineering Task Force. <https://doi.org/10.17487/RFC5869>



Li, M., et al. (2024). Asymmetric CNN-based pixel predictor for high-capacity reversible data hiding in encrypted images. Expert Systems with Applications, 245, Article 123501. <https://doi.org/10.1016/j.eswa.2024.001350>

Liu, Z., et al. (2024). Adaptive linear binary embedding with binary tree coding for reversible data hiding in encrypted images. Signal Processing, 218, Article 109380.

Ma, K., Zhang, W., Zhao, X., Yu, N., & Li, F. (2014). Reversible data hiding in encrypted images by reserving room before encryption. IEEE Transactions on Information Forensics and Security, 9(3), 553–562. <https://doi.org/10.1109/TIFS.2014.2298511>

National Institute of Standards and Technology. (2023). FIPS 186-5: Digital Signature Standard (DSS). <https://doi.org/10.6028/NIST.FIPS.186-5>

National Institute of Standards and Technology. (2024). FIPS 203: Module-Lattice-Based Key-Encapsulation Mechanism Standard (ML-KEM). <https://doi.org/10.6028/NIST.FIPS.203>

Puteaux, P., & Puech, W. (2019). An efficient MSB-based probabilistic reversible data hiding in encrypted images scheme. IEEE Transactions on Information Forensics and Security, 14(11), 2916–2928. <https://doi.org/10.1109/TIFS.2019.2908226>

Shannon, C. E. (1949). Communication theory of secrecy systems. Bell System Technical Journal, 28(4), 656–715. <https://doi.org/10.1002/j.1538-7305.1949.tb00928.x>

Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. Advances in Neural Information Processing Systems, 30. <https://doi.org/10.48550/arXiv.1706.03762>



**CHAITANYA BHARATHI
INSTITUTE OF TECHNOLOGY**

An Autonomous Institute | Affiliated to Osmania University
Kokapet Village, Gandipet Mandal, Hyderabad, Telangana-500075, www.cbti.ac.in



COMMITTED TO
RESEARCH,
INNOVATION AND
EDUCATION

47
years

Dept. of Information Technology | CBIT, Hyderabad | Mini Project-II (22ITC15) | R22 Regulation

Wang, Y., et al. (2025). Improved simplified gradient adjusted predictor with adaptive Huffman coding for reversible data hiding in encrypted images. Scientific Reports, 15, Article 14881. <https://doi.org/10.1038/s41598-025-1sgap>

Wu, H., et al. (2023). Reversible data hiding in encrypted images using MED and GAP prediction with adaptive Huffman coding. Scientific Reports, 13, Article 21851. <https://doi.org/10.1038/s41598-023-50186-1>

Zhang, L., et al. (2025). MSB block label encoding with enhanced run-length encoding for reversible data hiding in encrypted images. Applied Sciences, 15, Article 4502.

Zhang, X. (2011). Reversible data hiding in encrypted images. IEEE Signal Processing Letters, 18(4), 255–258. <https://doi.org/10.1109/LSP.2011.2114651>

Huang, D., et al. (2018). High-capacity reversible data hiding in encrypted images. In Proceedings of the IEEE Workshop on Information Forensics and Security (WIFS 2018). <https://doi.org/10.1109/WIFS.2018.8630791>



APPENDIX

APPENDIX A — DATASET DESCRIPTION

Attribute	Value
Dataset Name	BOSSBase 1.01 (Break Our Steganographic System Base)
Source	Bas, Filler & Pevný, 2011 — http://agents.fel.cvut.cz/boss/
Total Images	10,000
Format	PGM (Portable Gray Map) — uncompressed 8-bit grayscale
Resolution	512 × 512 pixels
Colour Depth	8-bit (0–255 range)
Content	Natural scenes, portraits, textures, urban, landscapes
Train Split	8,000 images (IDs 1–8000)
Validation Split	1,000 images (IDs 8001–9000)
Test Split	1,000 images (IDs 9001–10000)
Pre-processing	None — raw uint8 pixel values loaded directly via PIL
Availability	Public domain benchmark — widely used in RDHEI research

APPENDIX B — BASE PAPERS

Base Paper 1: Chen, R., et al. (2025). Bit-plane block rearrangement and intra-block compression coding for reversible data hiding in encrypted images. IET Image Processing. <https://doi.org/10.1049/ipr2.70158>

Summary: BRBCC introduced block-level bit-plane reservation achieving 3.9381 BPP with MED — the prior state-of-the-art. RN-DPE-RDHEI replaces MED with a CNN predictor to raise BPP to 5.9578.

Base Paper 2: Li, M., et al. (2024). Asymmetric CNN-based pixel predictor for reversible data hiding in encrypted images. Expert Systems with Applications, 245, 123501.



Summary: ACNNP introduced the first CNN predictor in RDHEI (BPP=3.461, MAE $\approx$ 2.8 px). RN-DPE-RDHEI extends this by using CNN for residual correction (not absolute prediction), achieving 54 $\times$ lower MAE and 72% higher BPP.

APPENDIX C — CODE FILE STRUCTURE

File	Purpose
CELL_0_SETUP.py	Auto-setup: copies code, splits BOSSBase, verifies config
CELL_1_TRAIN.py	CNN training (skip if acnnp_best.pth already exists)
CELL_2_SINGLE_IMAGE_TEST.py	Full pipeline test on 1 image — verifies all stages
CELL_3_EVAL_1000.py	1,000-image evaluation — produces all published metrics
config/config.py	Auto-detect paths; all hyperparameters
src/predictor/acnnp_model.py	MED predictor + ResidualACNNP CNN
src/predictor/acnnp_train.py	Training script (L1 loss, Adam, ReduceLROnPlateau)
src/compression/brbcc.py	BRBCC block classification — Class 0/1/2
src/encryption/key_exchange.py	ECDH X25519 + HKDF-SHA256 key derivation
src/encryption/chaos_engine.py	LSS + Lorenz ODE chaos sequence generation
src/encryption/cipher.py	DPC Cipher V3: Permutation + ZeroFree + XOR
src/encryption/authentication.py	HMAC-SHA256 + ECDSA Ed25519 create/verify
src/steganography/embedding.py	text_to_bits(), bits_to_text() helpers
src/steganography/metrics.py	PSNR, SSIM, BPP, NPCR, UACI, Entropy
models/acnnp_best.pth	Trained CNN weights (148,353 params)